Data Provenance & Methodology

Anvest Fintech Private Limited · Prepared for diligence · August 2026

1. Purpose and scope

This memo describes where Anvest's data comes from, how it is updated, what is verified today, and what is explicitly not yet verified. It is written for a technical or investment diligence reader and is deliberately specific about limitations — Anvest's product architecture treats an unstated gap as a worse failure than a disclosed one, and this document follows the same rule.

2. Data sources

- **End-of-day prices (bhavcopy):** the daily settlement file each exchange (NSE, BSE) publishes itself at market close. Anvest ingests this on a nightly schedule, currently 7:00 PM IST. This is not a real-time or licensed market-data feed — it is the exchanges' own published end-of-day record, which is why it carries no market-data licensing dependency.
- **Statutory filings and XBRL:** quarterly and annual financial results, annual reports, and corporate announcements as filed by companies with NSE/BSE under SEBI's listing obligations. These are public statutory disclosures, not a licensed third-party dataset.
- **OCR'd filings:** a meaningful share of small- and mid-cap filings are scanned PDFs rather than machine-readable XBRL. Anvest runs these through an OCR pipeline so they are searchable and citable on the same footing as machine-readable filings — this is the part of the corpus most competitors do not build.
- **Broker account data (opt-in):** where a user connects a brokerage account (read-only), holdings and transaction data personalize which filings and answers are surfaced. No trading or write access is requested.

3. Update cadence and the live-data roadmap

Today, the product runs on an end-of-day cycle: exchange bhavcopy and new filings are ingested in a nightly batch job. This is a deliberate, disclosed starting point, not a technical ceiling — the ingestion pipeline is designed to accept a streaming source in place of the nightly batch. The roadmap item is live quotes and live portfolio tracking through a broker's market-data and account API (e.g. Kite Connect), gated on the capital and hires this round funds; nothing about the current architecture needs to be rebuilt to support it.

4. Grounding architecture: how an answer is produced

Anvest's copilot does not calculate or estimate financial figures. Every numeric answer is read from the warehouse and returned with a citation to the specific filing and page it came from. Where coverage for a question does not exist, the product is designed to say so rather than produce a plausible-sounding number. This “grounded or silent” rule is the core reliability property of the product, and it is enforced at the retrieval layer rather than left to model behavior.

5. Coverage and disclosed gaps

Full-universe coverage currently spans 6,137 listed NSE/BSE symbols and roughly 584,000 parsed filings/announcements. Known, disclosed gaps include approximately 4,500 company-years for which the underlying filing was never made available by the exchange in a usable form — these are flagged in-product rather than silently omitted or backfilled with an estimate. Anvest's stated policy is to publish data-quality corrections as they are found rather than treat gap-discovery as reputationally costly; a discovered gap is closed by disclosure, not by concealment.

6. Derived metrics: verification status

Raw filing figures (revenue, profit, assets, equity, and similar reported line items) are read directly from the source filing and are the most directly verifiable layer — they can be checked against the filing PDF/XBRL tag itself, and Anvest's citation model is built to make that check easy for any user.

Derived ratios and metrics (ROE, ROCE, ROA, debt-to-equity, and similar computed figures) require additional logic — selecting the correct filing basis, handling consolidated vs. standalone figures, and applying sector-specific rules (for example, ROCE is not a meaningful figure for a lender and is deliberately withheld rather than computed). Anvest maintains an active internal data-quality program for this layer: automated identity checks (for example, that quarterly figures sum to the reported annual figure), cross-validation against independently recomputed values, and periodic comparison against an external public reference source on a sampled basis.

Current status, stated plainly: this verification program is ongoing, not complete. A corrected formula version has passed its internal engineering validation gates as of early August 2026, but has not yet been certified for production and has not yet been cut over from the version currently serving the product. In other words, ratio computation is an area of active, disclosed engineering work rather than a finished, externally audited system. Anvest will not represent derived-ratio accuracy more strongly than this until that work is complete and independently reviewable. Detailed technical status is available on request under a standard NDA/diligence arrangement.

7. Regulatory posture

Anvest does not currently hold SEBI Research Analyst (RA) registration; obtaining it is a funded priority for this round. Ahead of registration, the product enforces hard rules rather than relying on disclaimers: no buy/sell/hold recommendations, no target prices, and no personalized investment advice. SEBI's December 2024 AI-disclosure amendments require firms to disclose AI use in research output and hold the firm accountable for it — the “grounded or silent, always cited” architecture described in Section 4 was built with that accountability standard in mind, ahead of the registration requiring it.

8. Contact

Questions on any of the above, including access to the current internal data-quality report, can be directed to Harsh Bhardwaj, Founder, at harsh@anvest.in.