Data Provenance & Methodology

Anvest Fintech Private Limited · Prepared for diligence · October 2026

1. Purpose and scope

This memo describes where Anvest's data comes from, how it is updated, and how answers are grounded and verified. It is written for a technical or investment diligence reader. Anvest's product architecture treats an unstated gap as a worse failure than a disclosed one, and this document follows the same rule.

2. Data sources

- **End-of-day prices (bhavcopy):** the daily settlement file each exchange (NSE, BSE) publishes itself at market close. Anvest ingests this on a nightly schedule, currently 7:00 PM IST, with a morning catch-up run. Commercial use of exchange data is subject to each exchange's data policy; licence terms are being confirmed as part of the funded compliance work.
- **Statutory filings and XBRL:** quarterly and annual financial results, annual reports, and corporate announcements as filed by companies with NSE/BSE under SEBI's listing obligations. These are public statutory disclosures, not a licensed third-party dataset.
- **OCR'd filings:** a meaningful share of small- and mid-cap filings are scanned PDFs rather than machine-readable XBRL. Anvest runs these through an OCR pipeline so they are searchable and citable on the same footing as machine-readable filings — this is the part of the corpus most competitors do not build.
- **Broker account data (opt-in):** where a user connects a brokerage account (read-only), holdings and transaction data personalize which filings and answers are surfaced. No trading or write access is requested.

3. Update cadence and the live-data roadmap

Today the product runs on a daily cycle: exchange end-of-day files and new filings are ingested every evening, so the warehouse is current to the previous trading day. The next steps are in two layers. First, live filing alerts, which sit on the public exchange announcement feeds. Second, live quotes and portfolio tracking through each user's own broker connection (for example Kite Connect), with an exchange-licensed delayed feed for the public site. The ingestion pipeline is designed to accept these sources in place of, or alongside, the nightly batch.

4. Grounding architecture: how an answer is produced

Anvest's copilot does not calculate or estimate financial figures. Every numeric answer is read from the warehouse and returned with a citation to the specific filing and page it came from. Where coverage for a question does not exist, the product is designed to say so rather than produce a plausible-sounding number. This “grounded or silent” rule is the core reliability property of the product, and it is enforced at the retrieval layer rather than left to model behavior.

5. Coverage

The warehouse currently holds 6,137 listed NSE/BSE symbols, roughly 584,000 parsed filings, and 2.89M corporate announcements held in the warehouse, with scanned small-cap PDFs processed through OCR. Where a filing is unavailable in a usable form from the exchange, the product flags the gap in-product rather than estimating a figure. Data-quality corrections are published as they are found.

6. Verification of figures

Raw filing figures (revenue, profit, assets, equity, and similar reported line items) are read directly from the source filing and can be checked against the filing PDF or XBRL tag itself; Anvest's citation model makes that check one click for any user.

Derived ratios (ROE, ROCE, ROA, debt-to-equity and similar) go through an active data-quality program: automated identity checks (for example, that quarterly figures sum to the reported annual figure), cross-validation against independently

recomputed values, and sampled comparison against an external public reference. Sector rules are applied deliberately — ROCE, for instance, is withheld for lenders rather than computed. A corrected formula version has passed internal validation gates and is moving through certification to production. Detailed status is available on request under a standard diligence arrangement.

7. Regulatory posture

SEBI Research Analyst (RA) registration is a funded priority for this round. Until then, the product enforces hard rules rather than relying on disclaimers: no buy/sell/hold recommendations, no target prices, and no personalized investment advice. SEBI's December 2024 AI-disclosure amendments address disclosure of AI use in research output; the “grounded or silent, always cited” architecture in Section 4 was designed with that standard in mind. Counsel will confirm the product's regulatory position ahead of registration.

8. Contact

Questions on any of the above, including access to the current internal data-quality report, can be directed to Harsh Bhardwaj, Founder, at harsh@anvest.in.